VoiceOps 2026 — ZentroMart Support Agent

Project, run guide and evaluation report

This project implements section-aware policy retrieval, verified order lookup, human-ticket handoff, a Groq-hosted OpenAI-compatible conversational layer, and Gnani speech processing. A local browser interface supports text chat, evidence citations, session reset, microphone push-to-talk, and audio uploads.

Evaluation snapshot

Measure	Observed
Practice questions	52
Question decision baseline	52/52
Question gold-evidence coverage	47/47
ASR-text evidence coverage	46/47
Conversation turns	96
Conversation decision baseline	96/96
Conversation expected tool-call names	96/96

Step 1 vs Step 2

Step 1 prompt-only and generated-answer Step 2 quality were not measured in this evaluation. The table reports offline retrieval and decision baselines only; no generated-answer score is inferred from them.

Design and safeguards

Knowledge chunks retain section IDs and document metadata. Current policy, agent guide, help, FAQ and marketing precedence follows KB-00.1; superseded policy is excluded from ordinary retrieval. Order details remain gated behind the supplied mock verification tool, with one retry. Caller verification digits are redacted from stored conversation history, evidence files and ticket summaries. Support logs are not used as policy.

What went wrong and next steps

A live Groq JSON-mode check and an order-aware browser chat turn succeeded, including recovery from one provider JSON-formatting error. The full live practice evaluation did not complete because of provider rate limits and a connection error, so no Step 1 vs Step 2 scores are claimed here. The Gnani microphone/audio round trip still needs a consented recording test. Run `python run.py web` and open `http://127.0.0.1:5000`; then test Gnani audio and retry the full evaluation in smaller, rate-limited batches. Offline baselines are heuristics, not a substitute for human/native-language quality review.

Run the project

Install dependencies with `python -m pip install -r requirements.txt`; if `.env` does not already exist, copy `.env.example` to `.env` and add your own Groq and Gnani keys locally. Start the browser UI with `python run.py web`. For folder-based speech processing use `python run.py batch --input .\\my_audio --output .\\outputs`. See

README.md for setup, evaluation, and report commands.

Generated from the supplied synthetic dataset. Evaluation date: 2026-11-01.